

| | |
|-------------------------------------|---|
| snapshot | “xmin” and “xmax” are used to indicate what transactions are running at a specific time. When a transaction tries to read a table tuple, each tuple has a set of XIDs to indicate transactions which created and deleted the tuple. |
| xmin | when we insert rows, XID of inserting transaction is recorded at xmin field. |
| xmax | when we update rows, the old row is marked as "deleted" by writing updating transaction's XID to xmax field. At the same time the new updated row is created whose xmin field is marked with XID of the updating transaction. |
| gtm_List *gt_open_transactions | ./src/gtm/common/gtm_list.cpp ./src/include/gtm/gtm_list.h a dynamic list containing all open transactions at a given time. |
| GTM_TransactionInfo *gtm_txninfo | src/include/gtm/gtm_txn.h a fixed array with 8192 entries gt_lastslot : help locating new element from the gtm_txninfo. it keeps last used element in the array and search for an empty element for the new transaction from startslot = gt_lastslot+1 ; gt_lastfreed : keep the index of last freed gtm_txninfo array element. For a begin transaction call, if its value is valid, used it to get new element before checking gt_lastslot. Once used, gt_lastfreed is set to an invalid value and gets a valid value set only when the remove transaction call is executed. |
| LL_* | ./src/backend/libcomm/sctp_utils/sctp_uthash/utlsh.h |
| a write lock GTM_RWLockAcquire() | the duration of holding the lock would depend on (how soon an empty element can be located in the gtm_txninfo array) [ProcessingBeginTransactionCommand()] and (the number of entries in the open transactions list) [closedown]. |
| Two costly functions | GTM_GXIDToHandle() GTM_GetSnapshot() |

Src/include/

```
typedef struct GTM_Transactions
(
    uint32                gt_txn_count;
    GTM_States            gt_gtm_state;

    GTM_RWLock            gt_XidGenLock;

    /*
     * These fields are protected by XidGenLock
     */
    GlobalTransactionId gt_nextXid;        /* next XID to assign */
    GlobalTransactionId gt_backedUpXid;    /* backed up, restoration point */

    GlobalTransactionId gt_oldestXid;      /* cluster-wide minimum datfrozenxid */
    GlobalTransactionId gt_xidVacLimit;    /* start forcing autovacuum here */
    GlobalTransactionId gt_xidWarnLimit;   /* start complaining here */
    GlobalTransactionId gt_xidStopLimit;   /* refuse to advance nextXid beyond here */
    GlobalTransactionId gt_xidWrapLimit;   /* where the world ends */

    /*
     * These fields are protected by TransArrayLock.
     */
    GlobalTransactionId gt_latestCompletedXid; /* newest XID that has committed or
                                                * aborted */

    GlobalTransactionId gt_recent_global_xmin;

    int32                gt_lastslot;
    GTM_TransactionInfo gt_transactions_array[GTM_MAX_GLOBAL_TRANSACTIONS];
    gtm_List              *gt_open_transactions;

    GTM_RWLock            gt_TransArrayLock;
// JDU start
#ifdef GTM_LOCK_STATS
    int64                wrtCounts;
    timeval              startTm;
    timeval              accumTm;
    timeval              maxTm;
#endif
// JDU end
) GTM_Transactions;

extern GTM_Transactions GTMTransactions;
```

Src/include/

```
typedef struct GTM_TransactionInfo
{
    GTM_TransactionHandle    gti_handle;
    GTM_ThreadID             gti_thread_id;

    bool                     gti_in_use;
    GlobalTransactionId      gti_gxid;
    GTM_TransactionStates    gti_state;
    char                     *gti_coordname;
    GlobalTransactionId      gti_xmin;
    GTM_IsolationLevel       gti_isolevel;
    bool                     gti_readonly;
    GTMProxy_ConnID          gti_backend_id;
    char                     *nodelist; /* List of nodes prepared */
    char                     *gti_gid;

    GTM_SnapshotData         gti_current_snapshot;
    bool                     gti_snapshot_set;

    bool                     gti_vacuum;
    uint32                   gti_timeline;
    GTM_RWLock               gti_lock;
} GTM_TransactionInfo;
```

./src/gtm/main/gtm_txn.cpp

| | |
|---|--|
| ProcessBeginTransactionGetGXIDCommand() | |
| txn = GTM_BeginTransaction() -> GTM_BeginTransactionMulti() gxid = GTM_GetGlobalTransactionId(txn) -> GTM_GetGlobalTransactionIdMulti(txn, 1) -> GTM_HandleToTransactionInfo() | |
| | |

GTM_BeginTransactionMulti():

```
int ii, jj, startslot;^M

/**^M
 * We had no cached slots. Now find a free slot in the transaction array^M
 * and store the transaction info structure there^M
 */^M
startslot = GTMTransactions.gt_lastslot + 1;^M
if (startslot >= GTM_MAX_GLOBAL_TRANSACTIONS)^M
    startslot = 0;^M

for (ii = startslot, jj = 0;^M
     jj < GTM_MAX_GLOBAL_TRANSACTIONS;^M
     ii = (ii + 1) % GTM_MAX_GLOBAL_TRANSACTIONS, jj++)^M
{^M
    if (GTMTransactions.gt_transactions_array[ii].gti_in_use == false)^M
    {^M
        gtm_txninfo[kk] = &GTMTransactions.gt_transactions_array[ii];^M
        break;^M
    }^M

    if (ii == GTMTransactions.gt_lastslot)^M
    {^M
        GTM_RWLockRelease(&GTMTransactions.gt_TransArrayLock);^M
        ereport(ERROR,^M
                (ERANGE, errmsg("Max transaction limit reached")));^M
    }^M
}^M

init_GTM_TransactionInfo(gtm_txninfo[kk], coord_name, ii, isolevel[kk], connid[kk], readonly[kk]);^M

GTMTransactions.gt_lastslot = ii;^M

txns[kk] = ii;^M

/**^M
 * Add the structure to the global list of open transactions. We should^M
 * call add the element to the list in the context of TopMostMemoryContext^M
 * because the list is global and any memory allocation must outlive the^M
 * thread context^M
 */^M
GTMTransactions.gt_open_transactions = gtm_lappend(GTMTransactions.gt_open_transactions, gtm_txninfo[kk]);^M
```

GTM_GetGlobalTransactionIdMulti(txn, 1):

```

xid = GTMTransactions.gt_nextXid;^M
if (!GlobalTransactionIdIsValid(start_xid))^M
    start_xid = xid;^M

/*-----^M
 * Check to see if it's safe to assign another XID.  This protects against^M
 * catastrophic data loss due to XID wraparound.  The basic rules are:^M
 *^M
 * If we're past xidVacLimit, start trying to force autovacuum cycles.^M
 * If we're past xidWarnLimit, start issuing warnings.^M
 * If we're past xidStopLimit, refuse to execute transactions, unless^M
 * we are running in a standalone backend (which gives an escape hatch^M
 * to the DBA who somehow got past the earlier defenses).^M
 *^M
 * Test is coded to fall out as fast as possible during normal operation,^M
 * ie, when the vac limit is set and we haven't violated it.^M
 *-----^M
 */^M
if (GlobalTransactionIdFollowsOrEquals(xid, GTMTransactions.gt_xidVacLimit) &&^M
    GlobalTransactionIdIsValid(GTMTransactions.gt_xidVacLimit))^M
(^M
    if (GlobalTransactionIdFollowsOrEquals(xid, GTMTransactions.gt_xidStopLimit))^M
    (^M
        GTM_RWLockRelease(&GTMTransactions.gt_XidGenLock);^M
        ereport(ERROR,^M
            (ERANGE,^M
                errmsg("database is not accepting commands to avoid wraparound data loss in\
^M
            )^M
        else if (GlobalTransactionIdFollowsOrEquals(xid, GTMTransactions.gt_xidWarnLimit))^M
            ereport(WARNING,^M
                (errmsg("database must be vacuumed within %u transactions",^M
                    GTMTransactions.gt_xidWrapLimit - xid)));^M
    )^M

GlobalTransactionIdAdvance(GTMTransactions.gt_nextXid);^M
gtm_txninfo = GTM_HandleToTransactionInfo(handle[ii]);^M
if (gtm_txninfo == NULL)^M
(^M
    GTM_RWLockRelease(&GTMTransactions.gt_XidGenLock);^M
    ereport(ERROR, (EINVAL, errmsg("Invalid transaction handle")));^M
)^M

elog(DEBUG1, "Assigning new transaction ID = %d", xid);^M
gtm_txninfo->gti_gxid = xid;^M

```

GTM_HandleToTransactionInfo()

```

/**
 * Given the transaction handle, find the corresponding transaction info
 * structure
 */
/* Note: Since a transaction handle is just an index into the global array,
 * this function should be very quick. We should turn into an inline future for
 * fast path.
 */
GTM_TransactionInfo *
GTM_HandleToTransactionInfo(GTM_TransactionHandle handle)
{
    GTM_TransactionInfo *gtm_txninfo = NULL;

    if ((handle < 0) || (handle >= GTM_MAX_GLOBAL_TRANSACTIONS))
    {
        ereport(WARNING,
                (ERANGE, errmsg("Invalid transaction handle: %d", handle)));

        return NULL;
    }

    gtm_txninfo = &GTMTransactions.gt_transactions_array[handle];

    if (!gtm_txninfo->gti_in_use)
    {
        ereport(WARNING,
                (ERANGE, errmsg("Invalid transaction handle: %d, txn_info not in use", handle)));

        return NULL;
    }

    return gtm_txninfo;
}

```